
Angelika Langer
www.AngelikaLanger.com

The Art of Garbage Collection Tuning

objective


- discuss garbage collection algorithms in Sun/Oracle's JVM
- give brief overview of GC tuning strategies

agenda

- **generational GC**
- parallel GC
- concurrent GC
- "garbage first" (G1)
- GC tuning

three typical object lifetime areas


number of objects



generational GC

- key idea:
 - incorporate this typical object lifetime structure into GC architecture
- statically:
 - different heap areas for objects with different lifetime
- dynamically:
 - different GC algorithms for objects with different lifetime

static heap structure



different algorithms

- *mark-and-copy GC on young gen:*
collect objects with a short and short-to-medium lifetime
 - fast algorithm, but requires more space
 - enables efficient allocation afterwards
 - frequent and short pauses (*minor GC*)
- *mark-and-compact GC on old gen:*
collect objects which a medium-to-long and long lifetime
 - slow algorithm, but requires little space
 - avoids fragmentation and enables efficient allocation
 - rare and long pauses (*full GC*)

promotion

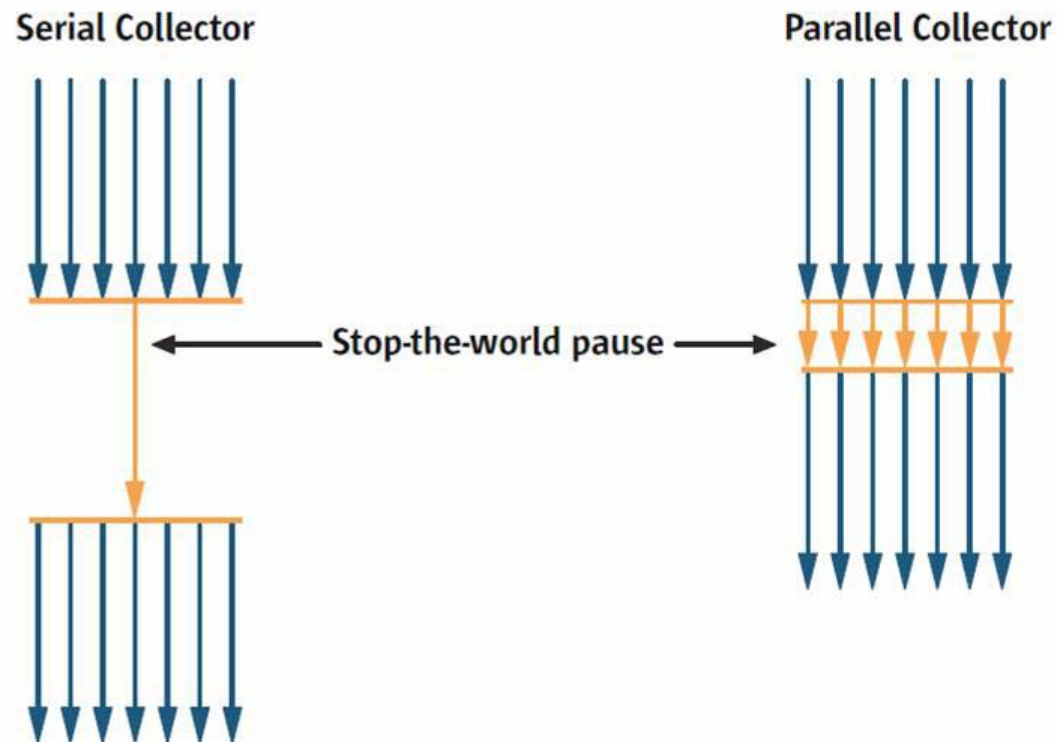
- aging and promotion
 - live objects are copied between survivor spaces and eventually to old generation
- how often live objects are copied between survivor spaces depends on ...
 - size of survivor space
 - number of live objects in eden and old survivor space
 - age threshold

agenda

- generational GC
- **parallel GC**
- concurrent GC
- "garbage first" (G1)
- GC tuning

multicore & multi-cpu architectures

- **parallel GC means:**
 - several GC threads + "stop-the-world"



parallel GC

- **parallel young GC** (since 1.4.1)
 - mark-sweep-copy
- **parallel old GC** (since 5.0_u6)
 - mark-sweep-compact
 - *mostly* parallel, i.e. has a serial phase

parallel young GC

- *mark phase (parallel)*
 - put all root pointers into a work queue
 - GC threads take tasks (i.e. root pointers) from work queue
 - GC threads put subsequent tasks (branches) into queue
 - *work stealing*: GC threads with empty queue "steal work" from another thread's queue (requires synchronization)
- *copy phase (parallel)*
 - challenge in parallel GC: many threads allocate objects in to-space
 - requires synchronization among GC threads
 - use *thread local allocation buffers* (GCLAB)

parallel old GC

- *marking phase (parallel)*
 - divide generation into fixed-sized regions => one GC thread per region
 - marks initial set of directly reachable live objects
 - keep information about size and location of live objects per region
- *summary phase (serial)*
 - determine *dense prefix*
 - point between densely and loosely populated part of generation
 - no objects are moved in dense prefix
 - loosely populated region is compacted
- *compaction phase (parallel)*
 - identify empty regions via summary data
 - parallel GC threads copy data into empty regions

agenda


- generational GC
- parallel GC
- **concurrent GC**
- "garbage first" (G1)
- GC tuning

concurrent old generation GC


- **concurrent GC means:**
 - no "stop-the-world"
 - one or several GC threads run concurrently with application threads
- **concurrent old GC** (since 1.4.1)
 - concurrent **m**ark-and-**s**weep GC algorithm (*CMS*)
 - goal: shorter pauses
 - runs *mark-and-sweep* => no compaction
 - works *mostly concurrent*, i.e. has stop-the-world phases

serial vs. concurrent old gc

Serial Mark-Sweep-Compact Collector



Concurrent Mark-Sweep Collector



concurrent old GC - details

- several phases
 - initial marking phase (serial)
 - marking phase (concurrent)
 - preclean phase (concurrent)
 - remarking phase (parallel)
 - sweep phase (concurrent)

concurrent old GC - details

- *initial marking (serial)*
 - identifies initial set of live objects
- *marking phase (concurrent)*
 - scans live objects
 - application modifies reference graph during marking
=> *not all live objects are guaranteed to be marked*
 - record changes for remark phase via *write barriers*
 - multiple parallel GC threads (since 6.0)

concurrent old GC - details

- *preclean (concurrent)*
 - concurrently performs part of remarking work
- *remarking (serial)*
 - finalizes marking by revisiting objects modified during marking
 - some dead objects may be marked as alive
 - => collected in next round (*floating garbage*)
 - multiple parallel GC threads (since 5.0)
- *sweep (concurrent)*
 - reclaim all dead objects
 - mark as alive all objects newly allocated by application
 - prevents them from getting swept out

CMS trace

- use `-verbose: gc` and `-XX: +PrintGCDetails` for details

| | |
|------------------------|--|
| non-concurrent mark | [GC [1 CMS-initial-mark: 49149K(49152K)] 52595K(63936K), 0.0002292 secs] |
| concurrent mark | [CMS-concurrent-mark: 0.004/0.004 secs] |
| concurrent preclean | [CMS-concurrent-preclean: 0.004/0.004 secs] [CMS-concurrent-abortable-preclean: 0.000/0.000 secs] |
| | [GC[YG occupancy: 3445 K (14784 K)] |
| non-concurrent re-mark | [Rescan (parallel) , 0.0001846 secs] |
| concurrent sweep | [weak refs processing, 0.0000026 secs] |
| concurrent reset | [1 CMS-remark: 49149K(49152K)] 52595K(63936K), 0.0071677 secs] [CMS-concurrent-sweep: 0.002/0.002 secs] [CMS-concurrent-reset: 0.000/0.000 secs] |

no mark and compact ...

CMS does not compact

- compacting cannot be done concurrently

downsides:

- fragmentation
 - requires larger heap sizes
- expensive memory allocation
 - no contiguous free space to allocate from
 - must maintain *free lists* = links to unallocated memory regions of a certain size
 - adverse affect on young GC (allocation in old gen happens during promotion)

fall back to serial GC

- CMS might not be efficient enough
 - to prevent low-memory situations
- CMS falls back to serial mark-sweep-compact
 - causes unpredictable long stop-the-world pauses

fall back to serial GC

concurrent GC —

```
[GC [1 CMS-initial-mark: 49149K(49152K)] 63642K(63936K), 0.0007233 secs]
[CMS-concurrent-mark: 0.004/0.004 secs]
[CMS-concurrent-preclean: 0.004/0.004 secs]
[CMS-concurrent-abortable-preclean: 0.000/0.000 secs]
[GC[YG occupancy: 14585 K (14784 K)]
[Rescan (parallel) , 0.0050833 secs]
[weak refs processing, 0.0000038 secs]
[1 CMS-remark: 49149K(49152K)] 63735K(63936K), 0.0051317 secs]
[CMS-concurrent-sweep: 0.002/0.002 secs]
[CMS-concurrent-reset: 0.000/0.000 secs]
```

serial GC —

```
[Full GC [CMS: 49149K->49149K(49152K), 0.0093272 secs] 63932K->63932K(63936K),
[CMS Perm : 1829K->1829K(12288K)], 0.0093618 secs]
```

out of memory —

```
[GC [1 CMS-initial-mark: 49149K(49152K)] 63933K(63936K), 0.0007206 secs]
java.lang.OutOfMemoryError
Heap
  par new generation   total 14784K, used 14784K
    eden space 13184K, 100% used
    from space 1600K, 100% used
    to   space 1600K,   0% used
  concurrent mark-sweep generation total 49152K, used 49150K
  concurrent-mark-sweep perm gen total 12288K, used 1834K
```

concurrent mark-and-sweep

- decreases old generation pauses
- at the expense of
 - slightly longer young generation pauses
 - some reduction in throughput
 - extra heap size requirements



agenda


- generational GC
- parallel GC
- concurrent GC
- **"garbage first" (G1)**
- GC tuning

garbage-first (G1) garbage collector

- available since Java 6 update 14 (experimental)
- features:
 - compacting
 - no fragmentation
 - more predictable pause times
 - no fall back to serial GC
 - ease-of-use regarding tuning
 - self adjustment; barely any options


general approach

- heap split into regions (+ perm)
 - 1 MByte each
- young region 
+ old region 
 - dynamically arranged
 - non-contiguous



young regions: collection

- copy live objects from young regions to survivor region(s)




young collection (details)

- coping of live objects = *evacuation pause*
 - stop-the-world, i.e. no concurrent execution of application
 - no good !!!
- but: evacuation is parallel
 - performed by multiple GC threads
 - good !!!
- parallel GC threads
 - GC operation is broken into independent tasks (work stealing):
 - determine live objects (marking stack)
 - copy live objects via GCLAB (similar to TLAB during allocation)

old regions: collection

- idea: collect regions with most garbage first
 - hence the name: "garbage-first"
- approach:
 - some regions may contain no live objects
 - very easy to collect, no coping at all
 - some regions may contain few live objects
 - live objects are copied (similar to young collection)
 - some regions may contain many live objects
 - regions not touched by GC

old regions: collection (cont.)



regions considered for GC evacuation

- *collection set* = regions considered for evacuation
- generational approach
 - sub-modes:
 - *fully young*: only young regions
 - *partially young*: young regions + old regions as pause time allows
 - GC switches mode dynamically
- which regions are put into the collection set ?
 - dynamically determined during program execution
 - based on a global marking that reflects a snapshot of the heap

note

- young and old regions have more similarities than before
- but still differences, i.e. it is generational GC
 - young regions:
 - where new objects are allocated
 - always evacuated
 - certain optimizations (e.g. no write barriers for remembered set update)
 - old regions:
 - only evacuated if time allows
 - only evacuated if full of garbage ("garbage-first")

benefits

- highly concurrent
 - most phases run concurrently with the application
 - some write barriers
 - some non-concurrent marking phases (similar to CMS)
 - even GC phases run concurrently
 - evacuation runs while global snapshot is marked
- highly parallel
 - multiple threads in almost all phases

benefits (cont.)

- fully self-adapting
 - just specify: max pause interval + max pause time
 - collection set is chosen to meet the goals
 - based on figures from various book keepings
 - e.g. previous evacuations, snapshot marking, write barriers

agenda

- generational GC
- parallel GC
- concurrent GC
- "garbage first" (G1)
- **GC tuning**

know your goals

- different applications require different GC behavior
 - no one-size-fits-all solution regarding GC and performance
- user aspects:
 - throughput
 - pauses
- engineering aspects:
 - footprint
 - scalability
 - promptness

profiling before you tune

- purpose
 - determine status quo
 - gather data for subsequent verification of successful tuning
- two sources
 - GC trace from JVM
 - profiling and monitoring tools

JVM options

- verbose: GC
- XX: +PrintGCDetails
 - switch on GC trace
 - details vary with different collectors
- XX: +PrintGCApplicationConcurrentTime
- XX: +PrintGCApplicationStoppedTime
 - measure the amount of time the applications runs between collection pauses and the length of the collection pauses

```
Application time: 0.5291524 seconds
[GC [DefNew: 3968K->64K(4032K), 0.0460948 secs] 7451K->
6186K(32704K), 0.0462350 secs]
Total time for which application threads were stopped:
0.0468229 seconds
```

JVM options

- XX: +PrintGCTimeStamps
 - enables calculation of total time, throughput, etc.
- Xloggc: <filename>
 - redirect GC trace to output file
- XX: +PrintTenuringDistribution
 - how often objects are copied between survivor spaces
- XX: +PrintHeapAtGC
 - prints description of heap before and after GC
 - produces massive amounts of output

heap snapshots

{Heap before GC invocations=1:

Heap

def new generation total 576K, used 561K [0x02ad0000, 0x02b70000, 0x02fb0000)

eden space 512K, 97% used [0x02ad0000, 0x02b4c7e8, 0x02b50000)

from space 64K, 100% used [0x02b60000, 0x02b70000, 0x02b70000)

to space 64K, 0% used [0x02b50000, 0x02b50000, 0x02b60000)

tenured generation total 1408K, used 172K [0x02fb0000, 0x03110000, 0x06ad0000)

the space 1408K, 12% used [0x02fb0000, 0x02fdb370, 0x02fdb400, 0x03110000)

compacting perm gen total 8192K, used 2433K [0x06ad0000, 0x072d0000, 0x0aad0000)

the space 8192K, 29% used [0x06ad0000, 0x06d305e8, 0x06d30600, 0x072d0000)

No shared spaces configured.

Heap after GC invocations=2:

Heap

def new generation total 576K, used 20K [0x02ad0000, 0x02b70000, 0x02fb0000)

eden space 512K, 0% used [0x02ad0000, 0x02ad0000, 0x02b50000)

from space 64K, 31% used [0x02b50000, 0x02b55020, 0x02b60000)

to space 64K, 0% used [0x02b60000, 0x02b60000, 0x02b70000)

tenured generation total 1408K, used 236K [0x02fb0000, 0x03110000, 0x06ad0000)

the space 1408K, 16% used [0x02fb0000, 0x02feb1b8, 0x02feb200, 0x03110000)

compacting perm gen total 8192K, used 2433K [0x06ad0000, 0x072d0000, 0x0aad0000)

the space 8192K, 29% used [0x06ad0000, 0x06d305e8, 0x06d30600, 0x072d0000)

No shared spaces configured.

}

GC trace analyzer - GCViewer

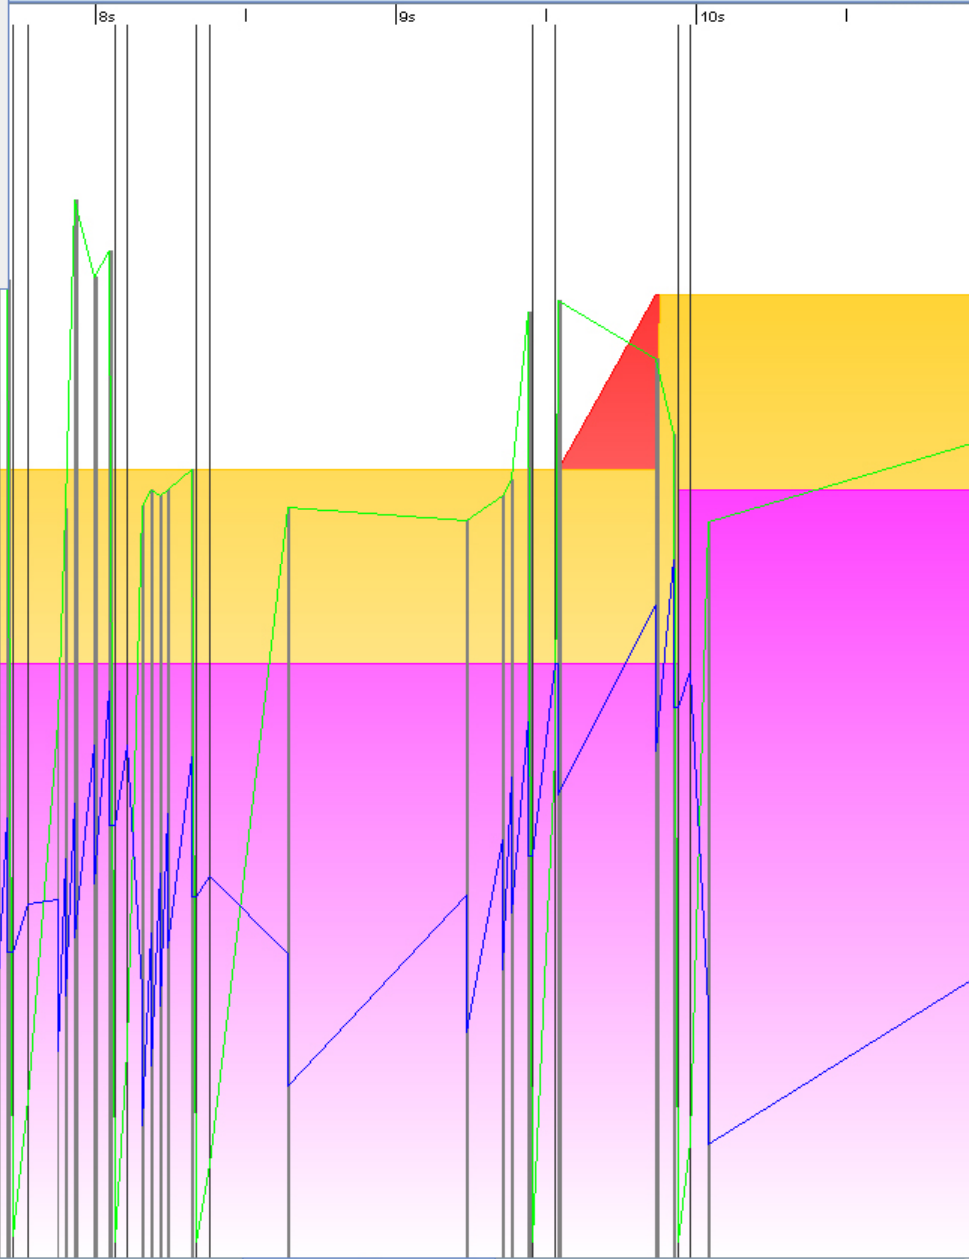
- GCViewer
 - freeware GC trace analyzer
 - until 2008 by Hendrik Schreiber at <http://www.tagtraum.com/gcviewer.html>
 - until 2008 by Jörg Wüthrich at <https://github.com/chewiebug/GCViewer>
- reads JVM's GC log file
 - post-mortem or periodically
- produces diagrams and metrics
 - throughput
 - pauses
 - footprint

- ☒ Daten-Tafel
- ☐ Antialias
- ☒ Vollständige GC Linien
- ☒ Inkrementelle GC Linien
- ☒ GC Dauer Linie
- ☒ GC Dauer Rechtecke
- ☒ Heapgröße
- ☒ Alte Generation
- ☒ Junge Generation
- ☒ Belegter Heap

25.000
24.000
23.000
22.000
21.000
20.000
19.000K
18.000K
17.000K
16.000K
15.000K
14.000K
13.000K
12.000K
11.000K
10.000K
9.000K
8.000K
7.000K
6.000K
5.000K
4.000K
3.000K
2.000K
1.000K
0K

0,017s
0,016s
0,015s
0,014s
0,013s
0,012s
0,011s
0,010s
0,009s
0,008s
0,007s
0,006s
0,005s
0,004s
0,003s
0,002s
0,001s
0,000s

PerformanceJava/Labs/Code/Solutions/08.02.GCTuning/src/gc.log.53.ParallelYoung.ConcurrentOld.txt



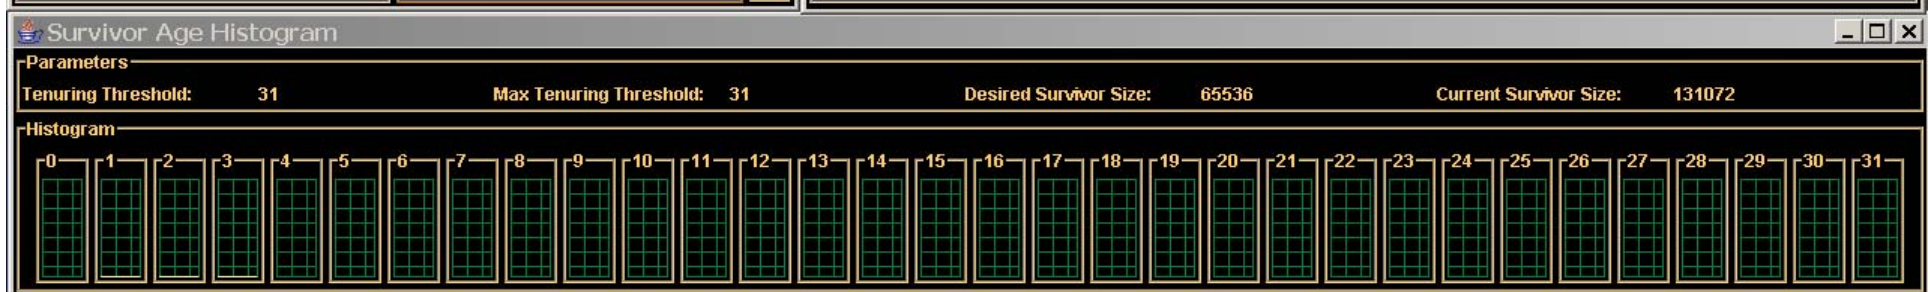


| Zusammenfassung | Speicher | Pause |
|-----------------------|----------|-------------------------------|
| Gesamtpausenzeit | | 0,87s |
| Summe vollst. GC | | 0,05s (5,4%) |
| Summe GC | | 0,82s (94,6%) |
| Kürzeste Pause | | 0,00027s |
| Längste Pause | | 0,02191s |
| Durchschn. Pause | | 0,01109s ($\sigma=0,00642$) |
| Durchschn. vollst. GC | | 0,00214s ($\sigma=0,00279$) |
| Durchschn. GC | | 0,01461s ($\sigma=0,00317$) |

| Zusammenfassung | Speicher | Pause |
|----------------------------------|----------|---------------------------------------|
| Gesamtspeicherverbrauch | | 24,898M |
| Durchschn. nach vollst. GC | | 11,168M ($\sigma=3.125,621K$) |
| Durchschn. nach GC | | 7.600,839K ($\sigma=3.489,794K$) |
| Insges. bereinigter Speicher | | 156,243M |
| Bereinigt von vollst. GC | | 0B (0,0%) |
| Bereinigt von GC | | 156,243M (100,0%) |
| Durchschn. bereinigt vollst. GC | | 0B/coll ($\sigma=0B$) |
| Durchschn. bereinigt von GC | | 2.857,018K/coll ($\sigma=221,988K$) |
| Durchschn. rel. Zuwachs nach VGC | | 376,692K/coll |
| Durchschn. rel. Zuwachs nach GC | | 980,906K/coll |
| Steigung nach vollst. GC | | 502,151K/s |
| Steigung nach GC | | 15,316M/s |

| Zusammenfassung | Speicher | Pause |
|------------------------------|----------|--------------|
| Gesamtspeicherverbrauch | | 24,898M |
| Insges. bereinigter Speicher | | 156,243M |
| Bereinigter Speicher/Min | | 509,378M/min |
| Gesamtlaufzeit | | 18s |
| Gesamtpausenzeit | | 0,87s |
| Durchsatz | | 95,3% |
| Vollst. GC Performance | | 0B/s |
| GC Performance | | 190,925M/s |

JVM monitor - VisualGC

- VisualGC
 - experimental utility (since JDK 1.4)
 - download from java.sun.com/performance/jvmsstat/visualgc.html
- integrated into VisualVM
 - download the VisualGC plugin (since JDK 6_u7)
- dynamically tracks and displays the heap
 - dynamic diagrams of all heap areas
 - no metrics at all



tuning for maximum throughput

- strategy #1: increase heap size
 - reduced overall need for GC
- strategy #2: let objects die in young generation
 - GC in old generation is more expensive than in young generation
 - prevent promotion of medium lifetime objects into old generation

let objects die in young generation

- increase young generation size
 - only limited by need for old generation size
- keep objects in survivor space
 - increase survivors space
 - raise occupancy threshold
 - raise age threshold
 - pro: prevents promotion of medium lifetime objects
 - con: needlessly copies around long lifetime objects
- use parallel young GC
 - increases throughput, if $\gg 2$ CPUs available

tuning for minimal pause time

- use parallel GC (parallel young and parallel compact)
 - reduces pause time, if $\gg 2$ CPUs available
- use concurrent GC (CMS)
 - pro: mostly concurrent
 - con: fragmentation + more expensive young GC
- try out "G1"
 - designed to limit pause time and frequency

tuning CMS


- strategy: avoid stop-the-world pauses
 - reduce duration of full GC
 - avoid full GC altogether

prevent fallback to stop-the-world GC

- increase heap size
 - defers the problems ("night time GC")
- start CMS early, i.e. lower occupancy threshold
 - reduces throughput because GC runs practically all the time
- increase young generation size
 - avoids fragmentation in the first place

tuning G1

- tuning G1 is different from classic GCs
 - generation sizes irrelevant
 - dynamically determined by G1 algorithms
 - still relevant: absolute memory size
 - grant as much memory as you can
- only 2 tuning parameters:
 - max pause + min interval



G1 tuning options

- MaxGCPauseMillis
 - upper limit for length of pause
 - what you demand from the GC
- GCPauseIntervalMillis
 - lower limit for length of interval in which GC pauses occur
 - how much GC activity you allow
 - short interval => many pauses in rapid succession
- defaults (might be too relaxed, for smaller apps)
 - GCPauseIntervalMillis = 500 ms
 - MaxGCPauseMillis = 200 ms

G1 tuning

- G1 "feels sluggish"
 - tuning goals are usually NOT met
- high variance compared to classic GCs
 - results differ even with identical tuning parameters
- G1 does not like overtuning
 - relaxed goal yields better results than ambitious goal

observations

- ambition is no good
 - raise pause time goal, i.e. demand shorter pause
 - (e.g. only 50 ms pause within 500 ms interval = 90% throughput)
 - result: G1 tries harder
 - make more pauses
 - often fails to reach the goal (pause time exceeds limit)
- relaxing is good
 - relax interval goal, i.e. allow more pauses
 - (e.g. 100 ms pause within 200 ms interval = only 50% throughput)
 - result: gives G1 more latitude and more flexibility
 - even pause times might decrease (without loss of throughput)
 - also avoids full GCs

wrap-up

- generational GC
 - split heap into generations
 - use different algorithms for each region
- young generation
 - mark-and copy (either serial or parallel)
 - many short stop-the-world pauses
 - needs survivor spaces

wrap-up

- old generation
 - mark-and-compact (either serial or parallel)
 - few gigantic stop-the-world pauses
 - no fragmentation
 - concurrent mark-and-sweep (CMS)
 - runs concurrently with the application
 - few short stop-the-world pauses (either serial or parallel)
 - falls back to mark-and-compact if needed
- "garbage first" (G1)
 - highly dynamic + very complex + hard to tune

wrap-up

- main tuning goals
 - throughput and pause times
- maximize throughput
 - let objects die in young generation
- minimize pauses times
 - avoid stop-the-world pauses

authors

Angelika Langer

Training & Consulting

Klaus Kreft

Performance Consultant, Germany

<http://www.AngelikaLanger.com>

garbage collection tuning

Q & A